

In Assignment 1, you have designed and implemented an efficient parallel algorithm to solve the **all pairwise computation** problem on shared-memory computing platforms and implemented the algorithm using OpenMP. In **Assignment 2**, you are asked to design and implement an efficient parallel algorithm to solve the same problem on a **distributed-memory machine**. We assume that in this distributed-memory machine there are a number of processing nodes and each processing node consists of a number of cores with shared memory. Thus, in your algorithm design and implementation, you need to consider data/tasks distribution and necessary communication among the processing nodes; you also need to consider how you will assign tasks to different cores in each processing node. You will implement your algorithm using **MPI** for data communication between processing nodes and **OpenMP** for thread management and synchronization in each processing node. The all pairwise computation problem is presented again below.

All Pairwise Computation is defined as performing computation (e.g., correlations) between every pair of the elements in a given dataset. In this assignment, we consider a set of N sequences, each being of length M , which are stored in a two-dimensional (2D) matrix of size N by M (i.e., N rows and M columns). We then calculate a dot product for every possible pair of sequences (row vectors) in the matrix.

A dot product of two sequences, or vectors is defined as

$$x \cdot y = \sum_{i=0}^{M-1} x_i y_i$$

If we allow a vector to pair with itself, then there are $N(N+1)/2$ such pairs in total for N vectors. Assuming N is equal to 5 for example, we then have the following 15 vector pairs:

(0, 0)(0, 1)(0, 2)(0, 3)(0, 4)

(1, 1)(1, 2)(1, 3)(1, 4)

(2, 2)(2, 3)(2, 4)

(3, 3)(3, 4)

(4, 4)

In the above dot product formula, (i, j) denotes the pair of sequences i and j . After the dot product computation for pair (i, j) , we will have a single value as the output. Let a_i and a_j be two sequences, namely the i^{th} and j^{th} row vectors in the 2D matrix, and $c_{ij} = a_i \cdot a_j$.

In your parallel algorithm design, you **must** consider how to (1) balance the workload, (2) minimize communication overhead, and (3) use the loop unrolling technique with the unrolling factor equal to 4 to improve the performance.

In your parallel algorithm implementation:

- Your program needs to ask for N (number of sequences) M (sequence length) and T (number of threads in each processing node) as user defined parameters;

- In your program process 0 first creates an input two-dimensional matrix of size N by M and then partitions and distributes the matrix evenly to other processes. As a result, the matrix is evenly divided into p (the number of created processes) blocks of rows and each process holds one block. It is not allowed for every process to store the whole input matrix during the computation.
- Since the dot products of (i, j) and (j, i) will give the same result, you must only compute the dot product of (i, j) for $i \leq j$ and store the results in the upper triangle of a 2D matrix as shown in the above example;
- After the parallel computation, your main program must conduct a self-check for correctness, i.e., perform a sequential computation using the same data set and then compare the result of this computation with the parallel computation result.

Your implementation must be in C and use MPI and OpenMP as described earlier in this assignment description.

You **must** write a **report**. The report must be concise, clear (3-6 A4 pages) and contain the following sections:

1. Problem definition and requirements
2. Parallel algorithm design and implementation
3. Testing
4. Discussion
5. Known issues in program
6. Manual (e.g. how to run the program, input and output)

Your assignment will be marked on the efficiency of your algorithm, program logic and readability, accuracy of results, and quality of your report.

You **MUST** attempt this assignment **individually**.

Submission Requirements

1. Your submission must be made by **11:59pm on Friday, 26 May, 2023 (Sydney time)**.
2. Create a **tar or zip file** that contains **your report**, **makefile** and **source files** (e.g., **.c** and **.h** files). **DO NOT INCLUDE ANY OBJECT OR BINARY FILES**.
3. Submit only one **.tar** or **.zip file**.

Failure to follow these submission requirements may lead to loss of marks.